





# Kubernetes Introduction

## Before Kubernetes




# Before Kubernetes


Traditional deployment



VM deployment



# Before Kubernetes – Traditional Deployment



# Before Kubernetes – Virtualization Deployment



Yes, but VM's have a few demerits as well

High cost of implementation

Hardware and software components are required by the VM's that increases the cost of implementation

Security risk

The chances of data breach while using virtualization is common

Availability issue

Assets are unavailable most of the time because the resources are being used by other applications

Limitations and restrictions

Not every application or server can work in a virtualization environment. Some may require hybrid machines to work with

Time-consuming

It is a very time-consuming process. This is because extra steps are followed to generate the desired results

# After Kubernetes

---

Yes, Kubernetes works  
very well and can be a  
good solution to  
overcome these issues



**kubernetes**

How about Kubernetes?  
It works with container  
deployment



# Virtual Machine vs Kubernetes

Major differences are:



Not secured



Security Risk

Not easily portable



Portability

Time is more



Time Consuming



**kubernetes**



Secured



Portable



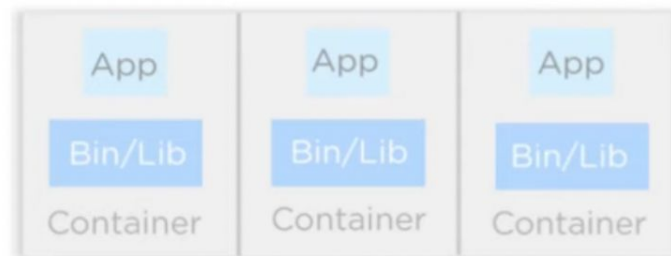
Time is less

# Virtual Machine vs Kubernetes



VM - They have less isolation when compared to Kubernetes, and hence, security risks are higher

Kubernetes - They have better isolation properties to share the OS among various applications



Container Runtime

OS

Hardware



# Kubernetes Era

---

Due to these reasons,  
Kubernetes is one of  
the best solution



**kubernetes**

Can we learn more  
about Kubernetes? It  
sounds fascinating to  
me



# What is Kubernetes?



**kubernetes**

In simple words, Kubernetes is an open-source platform used for maintaining and deploying a group of containers



# Benefits of Kubernetes

---

Portable and  
100% open-  
source

Workload  
scalability

High availability

Designed for  
deployment

Service  
discovery and  
load balancing


Kubernetes specifies how much CPU and RAM each container needs. Once resources are defined, managing the resources and making decisions for them becomes easy

Storage  
orchestration

Self healing

Automated  
rollouts and  
rollbacks

Automatic  
bin packing



**Worker Nodes**  
Host Application as Containers




## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers





**ETCD**  
CLUSTER

A stylized illustration of a Kubernetes cluster. On the left, a grey server rack contains a purple box labeled 'ETCD CLUSTER' with four circles inside. A vertical red line separates the control plane from the worker nodes. To the right of the red line is a large orange capsule-shaped node. A horizontal orange bar with a white truss bridge icon spans across the red line. The background is a blue sky with a blue body of water at the bottom.

**ETCD  
CLUSTER**

**kube-scheduler**

**kube-apiserver**



## Master

Manage, Plan, Schedule, Monitor  
Nodes



kube-apiserver



Master

Manage, Plan, Schedule, Monitor  
Nodes

kubelet

Controller-  
Manager

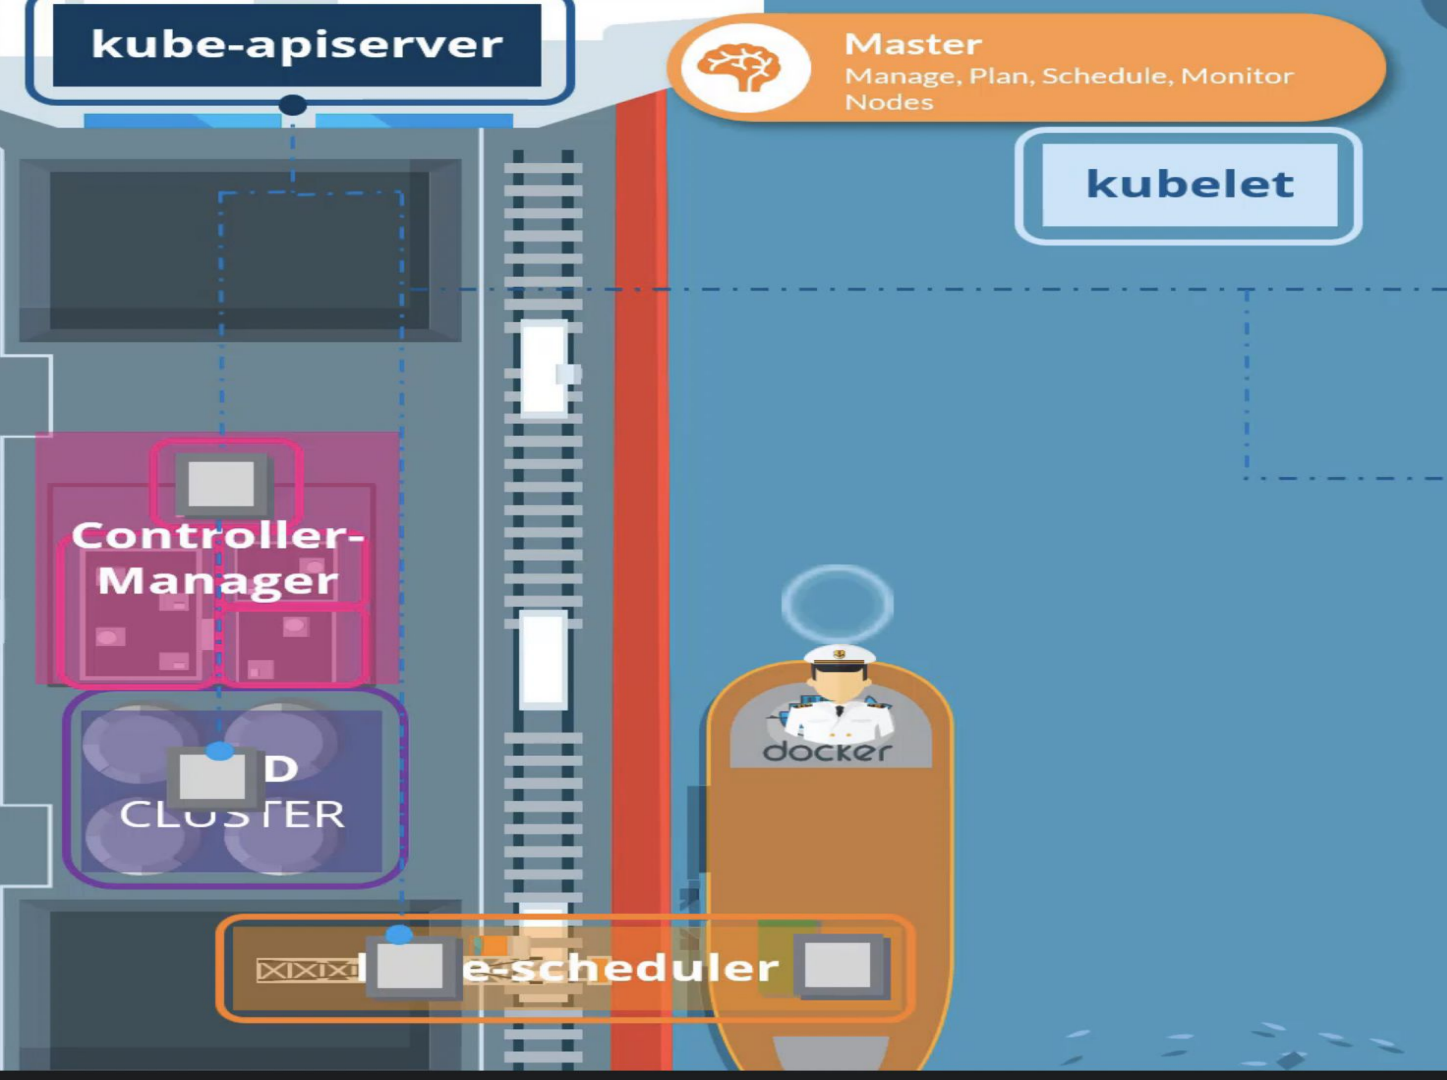
D  
CLUSTER



escheduler



docker



# Kubernetes Architecture



## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers

ETCD  
CLUSTER

kube-  
apiserver

Kube  
Controller  
Manager


kube-scheduler

kubelet

Kube-proxy

Container Runtime Engine

Run containers



kubelet

Kube-proxy

Container Runtime Engine


Run containers



Pods

# POD

---



replicaset-definition.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      type: front-end
```

pod-definition.yml

```
apiVersion: v1
kind: Pod
```

```
> kubectl create -f replicaset-definition.yml
```

```
replicaset "myapp-replicaset" created
```




```
> kubectl get replicaset
```


NAME	DESIRED	CURRENT	READY	AGE
myapp-replicaset	3	3	3	19s

# Labels and Selectors




# POD





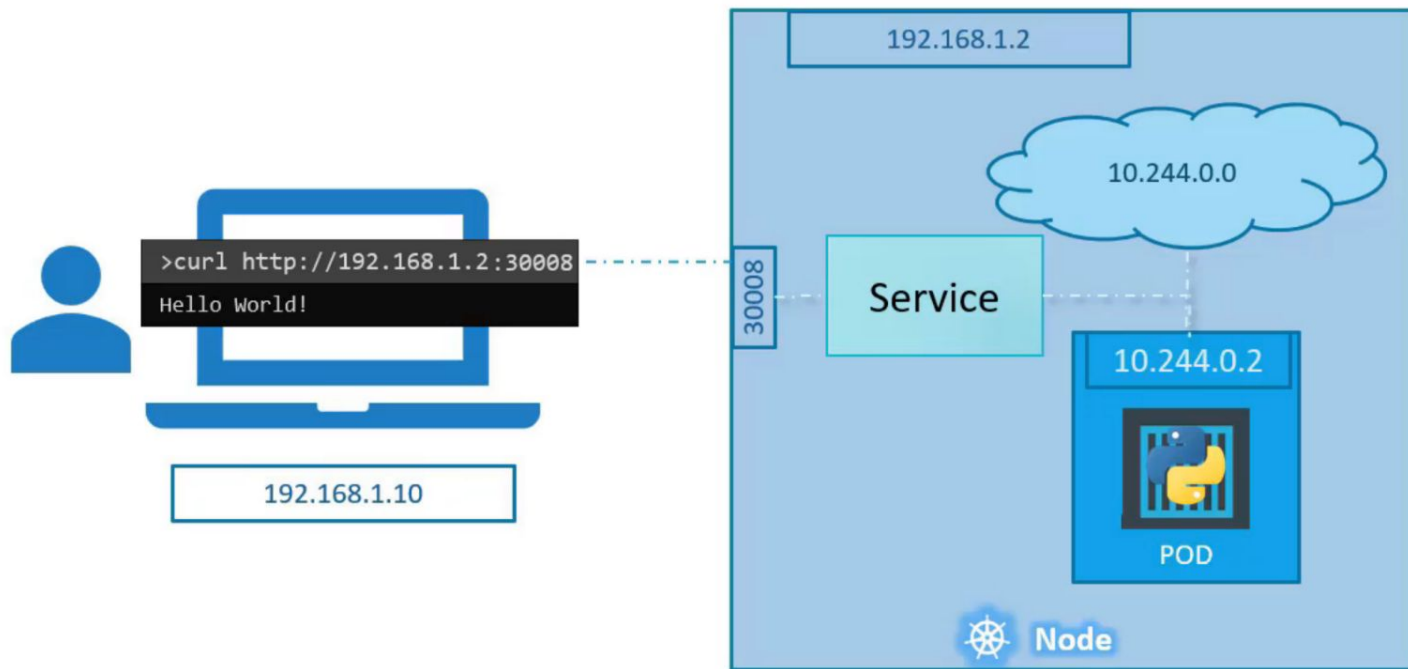
# Services

---





# Service




# Services Types


---



NodePort

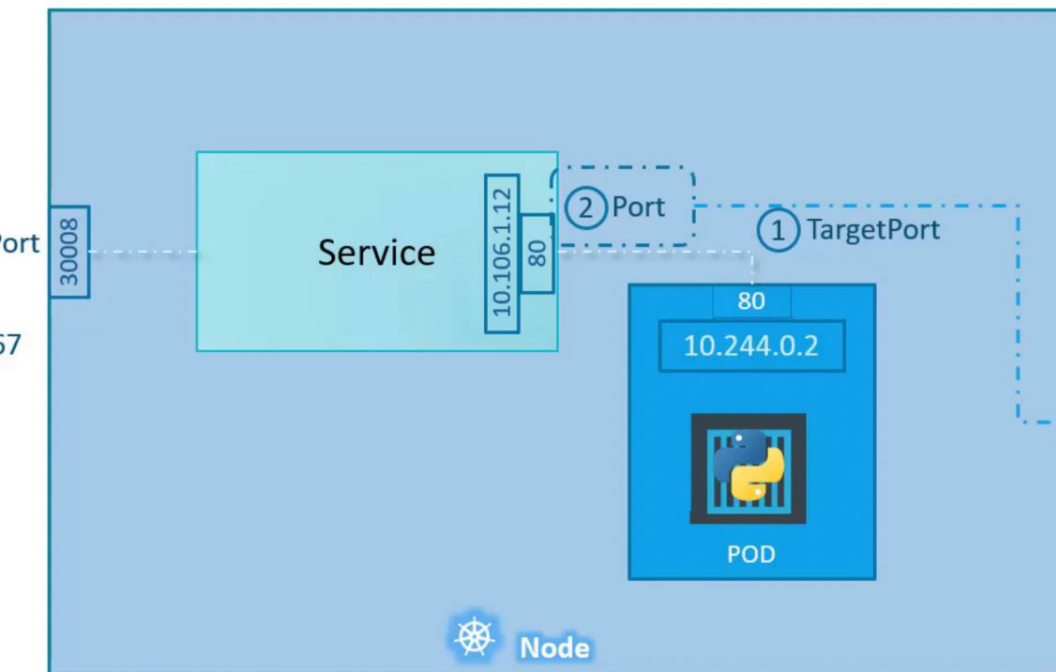


ClusterIP



LoadBalancer

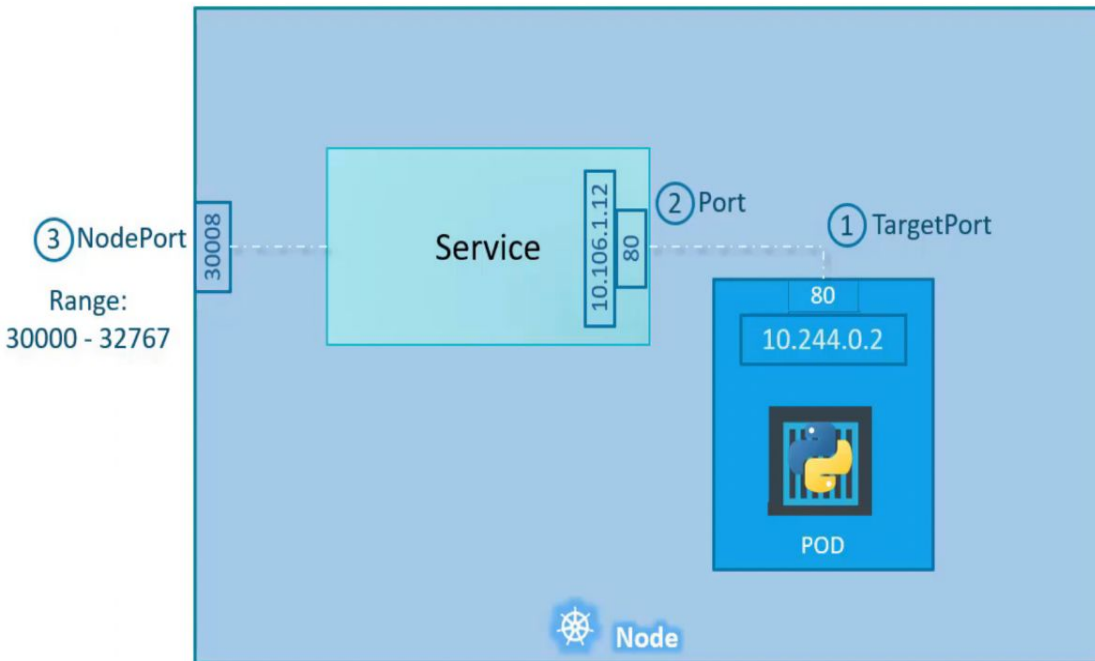
# Service - NodePort



service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
```




# Service - NodePort





service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

# Namespace - Isolation




# DNS




```
mysql.connect("db-service")
```

```
mysql.connect("db-serv
```


# Switch



```
> kubectl get pods --namespace=dev
```



```
> kubectl get pods
```



```
> kubectl get pods --namespace=prod
```

```
> kubectl config set-context $(kubectl config current-context) --namespace=dev
```

```
> kubectl get pods
```

```
> kubectl get pods --namespace=default
```

```
> kubectl get pods --namespace=prod
```

```
> kubectl config set-context $(kubectl config current-context) --namespace=prod
```

```
> kubectl get pods --namespace=dev
```

```
> kubectl get pods --namespace=default
```

```
> kubectl get pods
```

```
> kubectl get pods --all-namespaces
```

<https://minikube.sigs.k8s.io/docs/start/>




# | Daemon Sets – UseCase

Monitoring Solution

Logs Viewer

Daemon Sets



Step 1

Configure cluster

Step 2

Specify networking

Step 3

Configure logging

Step 4

Review and create

# Configure cluster

## Cluster configuration [Info](#)

**Name** - *Not editable after creation.*

Enter a unique name for this cluster.

javahome-eks

**Kubernetes version** [Info](#)

Select the Kubernetes version for this cluster.

1.20

**Cluster Service Role** [Info](#) - *Not editable after creation.*

Select the IAM Role to allow the Kubernetes control plane to manage AWS resources on your behalf.

To create a new role, go to the [IAM console](#).

Select role



## Secrets encryption [Info](#)

*Once enabled, secrets encryption cannot be modified or removed.*



**Enable envelope encryption of Kubernetes secrets using KMS**

Enable envelope encryption to provide an additional layer of encryption for your Kubernetes secrets.

Tags (0)

## Use case

Allow an AWS service like EC2, Lambda, or others to perform actions in this account.

---

### Common use cases


- ☐ **EC2**  
Allows EC2 instances to call AWS services on your behalf.
- ☐ **Lambda**  
Allows Lambda functions to call AWS services on your behalf.

### Use cases for other AWS services:

EKS 

- ☐ **EKS**  
Allows EKS to manage clusters on your behalf.
- ☒ **EKS - Cluster**  
Allows access to other AWS service resources that are required to operate clusters managed by EKS.
- ☐ **EKS - Nodegroup**  
Allow EKS to manage nodegroups on your behalf.
- ☐ **EKS - Fargate pod**  
Allows access to other AWS service resources that are required to run Amazon EKS pods on AWS Fargate.
- ☐ **EKS - Fargate profile**  
Allows EKS to run Fargate tasks.
- ☐ **EKS - Connector**  
Allows access to other AWS service resources that are required to connect to external clusters
- ☐ **EKS Local - Outpost**  
Allows Amazon EKS Local to call AWS services on your behalf.


# Deployment



When we need to update instance we need update  
docker instance seamlessly

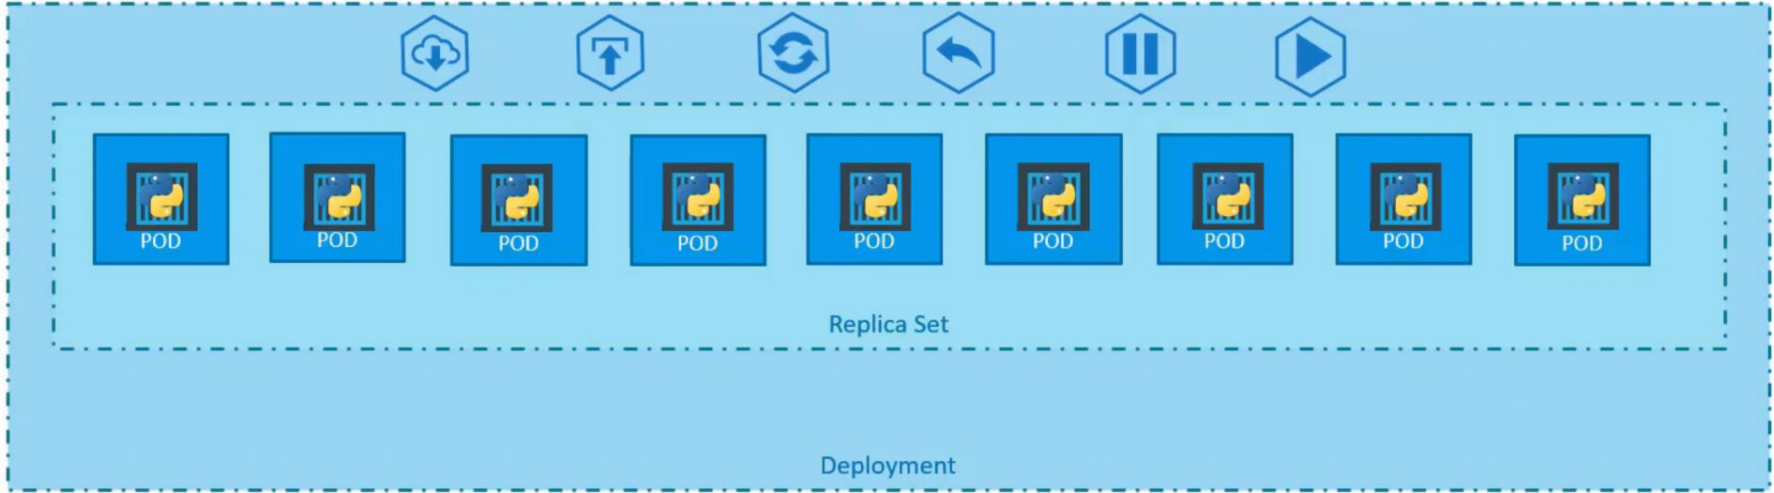
# Rolling update

We need to update the cluster with rolling update one after another we should be



Rolling update we should be able to undo changes which are recently carried out

Deployment



This will help in the doing upgrade seamlessly and undo changes resume changes as required

<https://docs.aws.amazon.com/eks/latest/userguide/sample-deployment.html>



<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>